

PROGRAMMING IN PYTHON I

Unit 08: Fast numerical computations and storing Python objects as files



Michael Widrich
Institute for Machine Learning

Copyright statement:

This material, no matter whether in printed or electronic form, may be used for personal and non-commercial educational use only. Any reproduction of this material, no matter whether as a whole or in parts, no matter whether in printed or in electronic form, requires explicit prior acceptance of the authors.

FAST NUMERICAL COMPUTATIONS IN PYTHON



Motivation

- We already learned that Python is an interpreted language
 - Very convenient to use
 - Slow, since optimization of the code is difficult at runtime
- We can use modules in Python that allow us to write fast code in Python
 - By providing optimized functions (e.g. NumPy, ...)
 - By providing tools for optimizing Python-like code (e.g. Numba, PyTorch, Tensorflow, ...)

NumPy

- NumPy is the go-to module for numerical computations in Python
- Provides a large range of functionalities for performing scientific computations and handling array data
 - These functions are typically highly optimized and implemented in C
- NumPy mainly deals with (multidimensional) array data based on the `numpy.ndarray` object
- Documentation/Tutorials:
<https://numpy.org/doc/stable/index.html>

Arrays in NumPy (1)

- We already heard about the simple linear array
 - Elements are stored as one block with contiguous addresses in memory
 - Elements are fast to access since we can quickly compute their addresses

Memory:

| | | | | | | | | | |
|-----|------|------|------|------|------|------|------|------|-----|
| ... | byte | byte | byte | byte | byte | byte | byte | byte | ... |
|-----|------|------|------|------|------|------|------|------|-----|

Address: ... 105 106 107 108 109 110 111 112 ...

Memory to store a 16-bit integer:

| | |
|------|------|
| byte | byte |
|------|------|

Storing 4 16-bit integers in memory:

| | | | | | | | | | |
|-----|------|------|------|------|------|------|------|------|-----|
| ... | byte | byte | byte | byte | byte | byte | byte | byte | ... |
|-----|------|------|------|------|------|------|------|------|-----|

... 105 106 107 108 109 110 111 112 ...

↑ ↑ ↑ ↑
Addresses of our integers

Arrays in NumPy (2)

- In Python, an element in a list is like a Python variable that is a reference to an object
 - Datatypes of objects are flexible
 - Operations on elements are slower/clumsy (need to determine type of object before usage)
- In Numpy, an element in an array is (usually) a bit-pattern that directly represents the stored value
 - The array holds the information about the datatype (encoding/decoding scheme for bits) used in array
 - Datatype of elements in array is fixed (but we can create new arrays with a different datatype)
 - All elements in an array have the same datatype
 - Operations on elements can be optimized better and are faster

Multidimensional arrays

- In Python lists, we already saw the concept of nested lists
 - Can be used to create 2-D or n-D arrays
 - Slow since we have to access the sub-lists to access our elements
- We can store n-D arrays as fast 1-D arrays
 - Done by NumPy in the background
 - Store n-D array in a **flat** manner
 - **Row-major** order: Consecutive elements of a row reside next to each other
 - **Column-major** order: Consecutive elements of a column reside next to each other

Multidimensional arrays: example (1)

- We want to store a 2-D array with 3 rows and 5 columns
 - 5 elements per row, 3 per column, 15 in total

| | | | | |
|----|----|----|----|----|
| 0 | 1 | 2 | 3 | 4 |
| 5 | 6 | 7 | 8 | 9 |
| 10 | 11 | 12 | 13 | 14 |

- We can create a 1-D array with 15 elements
- We can say that
 - the first 5 elements belong to the row in the first column
 - the next 5 elements belong to the row in the second column
 - the last 5 elements belong to the row in the third column
 - = row-major order

| | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|

Multidimensional arrays: example (2)

- We agreed on row-major order
- Now we want to access the element in the 4th column $c = 3$ and the 3rd row $r = 2$ (indices starting at 0 with $n_r = 5$ elements per row)

| | | | | |
|----|----|----|----|----|
| 0 | 1 | 2 | 3 | 4 |
| 5 | 6 | 7 | 8 | 9 |
| 10 | 11 | 12 | 13 | 14 |

- We can compute the index in the 1-D array via

$$n_r \cdot r + c = 5 \cdot 2 + 3 = 13$$

| | | | | | | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 |
|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|

Indexing in NumPy

- Accessing arrays in Numpy is similar to accessing lists in Python

- Index via integers:

```
my_array[my_index]
```

- Slicing is possible and fast (since elements are consecutively stored in memory):

```
my_array[:my_index]
```

- Numpy offers many more **fancy indexing** options

- Indexing multi-dimensional arrays directly:

```
my_array[my_row_index, my_col_index]
```

```
my_array[2, 4, 8, 5]
```

- Indexing using lists of indices, boolean index masks, ...

STORING PYTHON OBJECTS TO FILES



Pickle and dill: multi-purpose storage

■ pickle module

- ☐ Allows us to store and load many types of Python objects
- ☐ Stores data in binary files
- ☐ Not compressed (unless we compress the file using compression modules)
- ☐ Can handle many different Python objects

■ dill module

- ☐ Same interface as `pickle`
- ☐ Extends functionality of `pickle`
- ☐ Can store more types of Python objects
- ☐ Often used as `import dill as pickle`

H5py: storing large data

■ hdf5

- ☐ Container-format that allows for storing large data

■ h5py module

- ☐ Uses `hdf5` format to store Python objects in binary files
- ☐ Container-access via dictionary-like interface
- ☐ Array-access via NumPy-like indexing and slicing
- ☐ Supports different compression algorithms
- ☐ Accessing arrays can be done chunk-wise (e.g. if array does not fit in RAM at once)
- ☐ Types of store-able objects are limited