

# PROGRAMMING IN PYTHON II

## Model Evaluation for Classification Tasks



Michael Widrich  
Institute for Machine Learning

## Copyright statement:

This material, no matter whether in printed or electronic form, may be used for personal and non-commercial educational use only. Any reproduction of this material, no matter whether as a whole or in parts, no matter whether in printed or in electronic form, requires explicit prior acceptance of the authors.

# Outline

1. Motivation
2. Confusion matrix
3. True positive rate and true negative rate
4. Balanced Accuracy
5. Receiver Operating Characteristic (ROC)
6. Implementation
7. Excursion: Cross Validation

# MOTIVATION



# Motivation

- We have already learned about the **loss function**
  - Computes a **loss** given our model output and the target
  - The higher the loss, the farther away our output is to our target
  - We want to minimize the loss of our model during training


# Motivation

- We have already learned about the **loss function**
  - Computes a **loss** given our model output and the target
  - The higher the loss, the farther away our output is to our target
  - We want to minimize the loss of our model during training
- In this Unit we will learn about some common methods for evaluation of (binary) classifiers
  - Typically not used for training (require complete dataset to compute scores)
  - Well suited for comparison of classification performance with different goals

## Recap: Classification tasks and NNs

- Different classification tasks exist (see Unit 07)
- Binary classification task: Predict whether sample is in positive or negative class
- Typically done using a **sigmoid** activation function (aka logistic function) in the output layer
  - sigmoid function maps network output (“logit”) from range  $(-\infty, \infty)$  to range  $(0, 1)$
  - Activation after sigmoid can be interpreted as probability for belonging to positive class
  - Threshold to “classify” sample as positive or negative (threshold 0.5 is common):  
predicted\_class = positive if  $\text{sigmoid}(\text{logit}) \geq 0.5$ , otherwise  
predicted\_class = negative
  - Generalization to multiple classes via **softmax** function

# Sigmoid activation function





# CONFUSION MATRIX



# Confusion matrix

- The **confusion matrix** is a central piece for many evaluation methods
- Assume a binary classification task with samples belonging to either positive (P) or negative (N) class
- The confusion matrix illustrates the number of samples being
  - **True positive (TP)**: correctly classified as positive (aka “hit”)
  - **True negative (TN)**: correctly classified as negative (aka “correct rejection”)
  - **False positive (FP)**: wrongly classified as positive (aka “false alarm”, “Type I error”)
  - **False negative (FN)**: wrongly classified as negative (aka “miss”, “Type II error”)

# Confusion matrix for binary classification

		predicted class		
		positive	negative	total
true class	positive	TP	FN	P
	negative	FP	TN	N

# TRUE POSITIVE RATE AND TRUE NEGATIVE RATE



# True positive rate and true negative rate

## ■ True positive rate (TPR) (aka “sensitivity”, “recall”)

- ☐ Proportion of true positives that are correctly classified
- ☐  $TPR = \frac{TP}{P} = \frac{TP}{TP+FN}$
- ☐ E.g. % of diseased people that are correctly classified as being diseased

## ■ True negative rate (TNR) (aka “specificity”)

- ☐ Proportion of true negatives that are correctly classified
- ☐  $TNR = \frac{TN}{N} = \frac{TN}{TN+FP}$
- ☐ E.g. % of healthy people that are correctly classified as not being diseased

## ■ False negative rate (FNR) (aka “miss rate”)

- ☐  $FNR = \frac{FN}{P} = 1 - TPR$

## ■ False positive rate (FPR) (aka “fall-out”)

- ☐  $FPR = \frac{FP}{N} = 1 - TNR$

# BALANCED ACCURACY



# Balanced Accuracy

## ■ Accuracy (ACC) (aka “fraction correct (FC)”)

- Proportion of correctly classified samples
- $ACC = \frac{TP+TN}{P+N} = \frac{TP+TN}{TP+TN+FP+FN}$
- Simple measure but bad with imbalanced classes

## ■ Balanced accuracy (BACC)

- Normalizes TP and TN by the number of positive and negative samples
- $BACC = \frac{TPR+TNR}{2}$
- Suited for imbalanced classes

# RECEIVER OPERATING CHARACTERISTIC (ROC)





# Receiver Operating Characteristic (ROC)


## ■ Receiver operating characteristic (ROC) curve

- Shows TPR and FPR while varying classification threshold
- Classification performance w.r.t. FN and FP costs

## ■ Area under the receiver operating characteristic curve (AUC)

- Reduces ROC curve to single value (the area under the curve)
- Probability that a classifier will rank a randomly chosen positive instance higher than a randomly chosen negative one
- Ranges from 0 to 1, with 1 being perfect and 0.5 being random
- Suitable for imbalanced classes

# Receiver Operating Characteristic (ROC)



[Source: [https://scikit-learn.org/stable/auto\\_examples/model\\_selection/plot\\_roc.html](https://scikit-learn.org/stable/auto_examples/model_selection/plot_roc.html)]

# IMPLEMENTATION



# Implementation

- Many more evaluation methods exist, some of which with generalizations to tasks with multiple classes
- The `scikit-learn` (`sklearn.metrics`) module provides a vast selection of evaluation methods

`https://scikit-learn.org/stable/modules/classes.html#module-sklearn.metrics`

- For examples, see the code part of this Unit

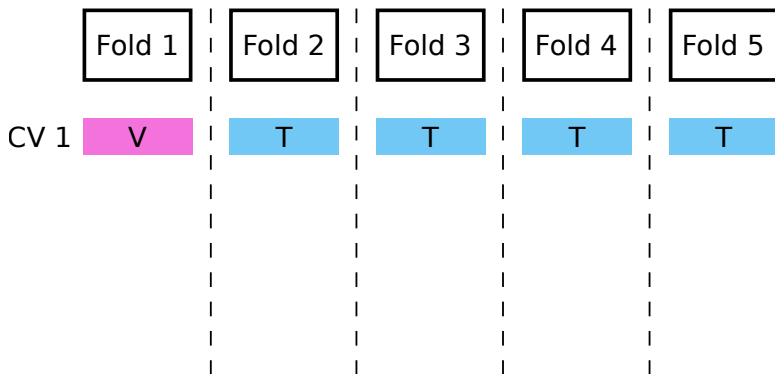
# EXCURSION: CROSS VALIDATION



## Excursion: Cross Validation

- To get an estimate for the model performance on future data, we have to use the test set to compute the scores
- For small datasets, the requirement that training and test set must not overlap is painful
- Solution: **Cross Validation (CV)**
  - Split dataset into  $n$  disjoint folds
  - Use  $n - 1$  folds as training set, left-out fold as test set
  - Train  $n$  times, every time leaving out a different fold as test set
  - Average over  $n$  estimated risks on test sets to get better estimate of generalization capability

# Cross Validation (1)



5-fold Cross Validation  
T: Training set; V: Test set

# Cross Validation (1)

	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5
CV 1	V	T	T	T	T
CV 2	T	V	T	T	T

5-fold Cross Validation  
T: Training set; V: Test set



# Cross Validation (1)

	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5
CV 1	V	T	T	T	T
CV 2	T	V	T	T	T
CV 3	T	T	V	T	T
CV 4	T	T	T	V	T
CV 5	T	T	T	T	V

5-fold Cross Validation  
T: Training set; V: Test set

## Cross Validation (2)

### ■ Nested Cross Validation

- We can apply another (inner) CV procedure within each training-set of the original (outer) CV
- allows for evaluation of model selection procedure

## Cross Validation (2)

### ■ Nested Cross Validation

- We can apply another (inner) CV procedure within each training-set of the original (outer) CV

→ allows for evaluation of model selection procedure

### ■ Getting a risk estimate on selected model:

1. Apply cross validation on training set (withhold test set)
2. Use test set to estimate risk for the model selected via CV

## Cross Validation (2)

- Nested Cross Validation

- We can apply another (inner) CV procedure within each training-set of the original (outer) CV

- allows for evaluation of model selection procedure

- Getting a risk estimate on selected model:

1. Apply cross validation on training set (withhold test set)
2. Use test set to estimate risk for the model selected via CV

- In practice, the found model is often trained further or re-trained on complete dataset for best performance